

NAIST Lecture (Nov. 12, 2025)

Shohei Higashiyama

Affiliate Assistant Professor@NAIST NLP Lab

NLP 2: Lexical Analysis

—Word Segmentation and Part-Of-Speech Tagging

Table of Contents:

1. Fundamental Concepts and Background
2. Some Methods for Lexical Analysis Tasks
3. Tokenization in the Neural NLP Era
4. The Usage of Foundational Lexical Analysis in Current NLP

Announcements for This Session

● Materials

- Lecture materials are stored in the following folder:
 - <https://drive.google.com/drive/folders/18gbP6ZwMhYC1QJwO5th29cw7tYsHR97f?usp=sharing>
- You can find this URL from the syllabus (<https://edu-portal.naist.jp/>)

● Assignments

- This session requires you to submit a report (some simple quizzes related to lexical analysis and tokenization).
- Assignment details will be announced at the end of the lecture.

[Slido link for Q&A](#)



● Questions

- If you have any questions, please submit them through Slido:
 - <https://app.sli.do/event/gdPg4DdX4RN1jaiNBhv6fh>


Part 1: Fundamental Concepts and Background

Task

● What is a *task* in NLP?

- A problem that a system or model needs to solve.
- Defined by its input and output.

Example: Recognizing Textual Entailment (RTE)



Example: Machine Translation (MT)

Text in the source language

体調はどうですか？

taichō wa dō desuka



Text in the target language

How are you feeling?

Representative Tasks in Lexical Analysis

- Tokenization
- Word Segmentation
- Part-of-Speech Tagging
- Morphological Analysis

[Notes]

- The term *Lexical Analysis* is sometimes used to encompass various lexical-level linguistic analysis tasks, but it is not as well-established as individual task names.
- This lecture regard tokenization as a lexical analysis-related task for convenience, but it not usually considered as such, as it does not necessarily identify linguistic units.

Tokenization and Word Segmentation

● Token/Tokenization

- A *token* refers to a unit of text for processing in NLP.
- Classical tokenizers (like NLTK) split a sentence into words and punctuation marks as tokens.



● Word Segmentation (=a type of tokenization)

- A task of dividing an *unsegmented sentence* into words.
- Typically performed for *unsegmented languages* w/o spaces, such as Japanese and Chinese.



Part-of-Speech Tagging

● Part-of-Speech (POS)

- A grammatical category of a word, such as *noun* and *verb*, which indicates how it functions in a sentence.
- Helps an NLP system understand the roles of words within a sentence.

● POS Tagging

- A task of assigning POS tags for words in a sentence.

Sequence of words

["Mr.", "Smith", "is", "n't", "worried."]



Sequence of POS tags

["PROPN", "PROPN", "AUX", "PART", "ADJ", "PUNCT"]

Universal POS tags (Universal Dependencies Project)

<u>ADJ</u> : adjective	<u>PART</u> : particle
<u>ADP</u> : adposition	<u>PRON</u> : pronoun
<u>ADV</u> : adverb	<u>PROPN</u> : proper noun
<u>AUX</u> : auxiliary	<u>PUNCT</u> : punctuation
<u>CCONJ</u> : coordinating conjunction	<u>SCONJ</u> : subordinating conjunction
<u>DET</u> : determiner	<u>SYM</u> : symbol
<u>INTJ</u> : interjection	<u>VERB</u> : verb
<u>NOUN</u> : noun	<u>X</u> : other
<u>NUM</u> : numeral	

Morphological Analysis


- Morpheme: The smallest linguistic unit that carries meaning.

Word Morphemes
runs ➤ run + s
Stem Suffix

- Definition of “Morphological Analysis”

- Broad sense: A general term for analysis-focused morphological learning tasks
- Narrow sense: A specific term refers to the combination of *lemmatization* and *morphological tagging*

* The *lemma* and *root* are also “run.”



Cited From [\[Liu '21\]](#)

* MSD: Morphosyntactic description

* MSD tags can be regarded as fine-grained POS tags.

Japanese Morphological Analysis

● Definition

- Typically refers to a complex sentence-level task, involving word segmentation, POS tagging, and lemmatization (inflection processing).


Input sentence		Output
昨日は楽しかった。	▶	Word: 昨日 は 楽しかつ た
kinō wa tanoshikatta		Lemma: 昨日 は 楽しい た
I had a good time yesterday.		POS tag: Noun Particle Verb Suffix

[Notes]

- Each “word” token actually correspond to a word, morpheme, or an intermediate unit. The segmentation granularity depends on the segmentation criterion (POS tag system).
- Thus, in Japanese NLP, *morphemes* and *words* are usually not strictly distinguished and are often used interchangeably.

Pipelined Task Flow in Traditional NLP


- In traditional NLP, lexical analysis tasks were essential preprocessing steps.



* Each application-oriented task does not necessarily require all the preceding lower-level tasks.

End-to-End Task Flow in Current NLP

- In current neural NLP,
tokenization is usually the only mandatory prerequisite step for a target task.



Other Fundamental Concepts

● Vocabulary

- A set of words/tokens collected based on specific criteria.
- An NLP model has a vocabulary of tokens that the model can process.

● Type vs. Token

- Word type: Unique form of a word in a text/vocabulary.
- Word token: Each instance of a word appeared in a text.

This sentence has
5 word tokens & 4 word types.

“a person has a pen”

● Ambiguity

- The possibility of multiple interpretations.
- *Word segmentation ambiguity* can impact the performance of downstream tasks.

米原 発熱 海 行き	→	For Maibara Heat Sea
米原発熱海行き	→	U.S. Nuclear Power Plant to Atami
Maibara-hatsu Atami-yuki	→	From Maibara to Atami Appropriate

- These are machine translation outputs that I obtained on a previous day.
- 米原 (Maibara) and 熱海 (Atami) are place names, so the third one is natural.

* NLP faces many challenges due to various kinds of ambiguities in languages!

Other Fundamental Concepts

● Vocabulary

- A set of words/tokens collected based on specific criteria.
- An NLP model has a vocabulary of tokens that the model can process.

● Type vs. Token

- Word type: Unique form of a word in a text/vocabulary.
- Word token: Each instance of a word appeared in a text.

This sentence has
5 word tokens & 4 word types.

“a person has a pen”

● Ambiguity

- The possibility of multiple interpretations.
- *Word segmentation ambiguity* can impact the performance of downstream tasks.

米原 発熱 海 行き	→	For Maibara Heat Sea
米原発熱海行き	→	U.S. Nuclear Power Plant to Atami
Maibara-hatsu Atami-yuki	→	From Maibara to Atami Appropriate

- These are machine translation outputs that I obtained on a previous day.
- 米原 (Maibara) and 熱海 (Atami) are place names, so the third one is natural.

* NLP faces many challenges due to various kinds of ambiguities in languages!

Summary of Part 1

- We have looked at:
 - Tokenization as a prerequisite step.
 - Definitions of lexical analysis tasks: word segmentation, POS tagging, and morphological analysis.
 - Fundamental concepts: token, vocabulary, and word tokens/types.
 - The role of tokenization and lexical analysis in the traditional/current NLP task flows.

Somewhat Focusing on
Word Segmentation


Part 2: Some Methods for Lexical Analysis Tasks

Progression of Methods for Word Segmentation and POS Tagging

- From around 2000 to the mid-2010s:
 - Statistical machine learning models were primarily used.

cf. [https://aclweb.org/aclwiki/POS_Tagging_\(State_of_the_art\)](https://aclweb.org/aclwiki/POS_Tagging_(State_of_the_art))

[Liu+ '23] Survey on Chinese Word Segmentation




A lattice structure for the HMM-based POS tagger [Lee+ '00]

- From the mid-2010s to the early 2020s:

- Neural network models were actively developed and have achieved substantial performance improvements (particularly in Chinese word segmentation).

Note: Statistical models, such as [MeCab](#) and [Jieba](#), remain popular as practical tools.



BiLSTM-CRF Tagger [Huang+ '15]
(Illustrated outputs are labels for an NER task.)

Statistical Methods for Japanese Morphological Analysis (JMA)


● Characteristics of statistical methods

- Typical methods are *lattice-based* approaches relying on dictionaries.
- Computationally efficient and fast.
- Available off-the-shelf models can achieve high accuracy for well-formed text (e.g., news).
 - * It is because such models are trained with texts like news articles.

● MeCab [[Kudo+ '04](#)]

- One of the de facto standard tools for JMA.
- Based on Conditional Random Fields.

1. Construct the *lattice* for each input sentence while referring to the MA dictionary.
2. Search the best path using the *Viterbi algorithm* based on the trained model parameters.



A **word lattice** is a graph in which nodes represent words, and edges indicate that the two nodes can be connected.

Neural Methods for Word Segmentation

● Characteristics of neural methods

- Often treat word segmentation as *sequence labeling*.
- Language-independent and easily extensible to a multi-task model.
- Can obtain benefit from powerful *pretrained language models* like BERT.

Sequence labeling:
a type of a task that predicts
a label sequence for an input
token sequence

Input character tokens:

[奈, 良, に, は, 鹿, が, い, る]



Label sequence to be predicted:

[B, E, B, B, B, B, B, E]

奈良|には|鹿|が|いる

Nara ni-wa shika-ga iru
There are deer in Nara.


B: Beginning of a word
E: End of a word

* Popular tag schemas:
BE (=0/1), BIE, BIES, etc.

* Many NLP tasks other than word segmentation can also be formulated as sequence labeling.

Neural Methods for Word Segmentation

- A character-word hybrid model [Higashiyama+ '19]
 - Relied a BiLSTM-CRF, which was the de facto standard architecture for neural sequence labeling (before the BERT era).
 - Achieved state-of-the-art accuracy for both Japanese and Chinese word segmentation.
 - By changing the training data, the same model can be used for different languages.



Neural Methods for Word Segmentation

● Model comparison on various Japanese texts

- MeCab achieved high accuracy for GEN (general) domain.
 - The model was trained on GEN domain data.
- BERT achieved high accuracy for many domains.
 - The model was pretrained on Japanese Wikipedia texts and fine-tuned on GEN domain data.

* BERT [Devlin+ '19] is a powerful neural NLP model and a prominent example of a masked language model (MLM).

Neural methods (based on pretrained language models) have the advantage of robustness to domain shift.

ENE to EMR: Scientific documents.
DIE to PRM: Government documents.
TBK to VRS: Other documents.

Accuracy (F1 scores) of Japanese word segmenters on various domains ("dom.") [Higashiyama+ '22]

Dom.	Unknown Tok/Type Ratio	MeCab		BL-LWP		BERT	
		D_s	Seg POS	D_s, D_t, U_t	Seg POS	– Seg	POS
GEN	2.7 / 16.1	99.6	99.0	98.9	98.3	99.4	99.1
ENE	2.5 / 15.4	99.3	98.9	99.6	99.2	99.7	99.4
TRA	3.0 / 18.2	98.8	98.4	99.4	98.9	99.6	99.2
ENV	3.2 / 15.1	98.8	98.1	99.3	98.7	99.5	99.2
MAN	3.3 / 19.5	98.6	98.2	99.4	99.0	99.6	99.3
CON	3.5 / 19.5	98.9	98.1	99.2	98.6	99.5	99.1
AGR	4.5 / 21.0	98.5	98.0	99.0	98.4	99.4	99.0
THM	4.5 / 24.0	98.4	97.7	99.1	98.3	99.4	98.8
INF	4.7 / 22.6	97.9	97.5	99.1	98.5	99.5	99.1
MEC	5.0 / 25.3	98.4	97.8	99.3	98.7	99.5	99.1
NUC	5.3 / 20.2	98.1	97.3	98.9	98.0	99.4	98.9
CHE-I	5.5 / 23.7	97.9	97.3	99.0	98.3	99.5	99.0
ETH	5.5 / 24.5	98.5	97.8	99.3	98.4	99.4	98.8
MED	5.6 / 27.0	97.1	96.6	99.1	98.6	99.5	99.1
SYS	5.6 / 24.8	98.4	97.7	98.9	98.0	99.4	98.7
ELC	5.8 / 29.4	97.4	97.0	99.0	98.5	99.5	99.1
PAT	6.0 / 26.8	97.0	96.8	99.1	98.7	99.4	99.2
CHE-E	6.1 / 23.7	97.9	97.0	99.0	98.0	99.2	98.7
MIN	6.6 / 22.6	98.0	97.4	98.8	98.1	99.0	98.6
BIO	6.7 / 30.2	96.7	96.0	98.8	98.0	99.3	98.7
PHY	7.5 / 29.6	97.1	96.4	98.5	97.7	99.2	98.8
CHE-B	8.1 / 35.4	97.0	96.1	98.5	97.4	99.1	98.4
EMR	11.2 / 30.2	95.4	91.9	95.6	92.5	97.1	94.0
DIE	0.9 / 7.5	98.0	97.6	97.7	97.0	97.9	97.4
LAW	2.1 / 11.0	97.4	97.0	97.6	97.4	97.9	97.8
PRM	2.8 / 11.1	98.7	98.1	97.7	96.8	98.3	97.8
TBK	4.4 / 19.0	99.0	97.0	97.7	95.5	98.7	96.8
VRS	19.7 / 47.4	87.3	82.3	81.8	75.1	87.1	83.0

Practical Systems for Japanese Word Segmentation

- Statistical methods, particularly MeCab, are still widely used for high processing speed.
- Juman++ V2 [\[Tolmachev+ '20\]](#) utilizes a recurrent neural network (RNN) language model to achieve high accuracy while maintaining practical processing speed.


Processing speed for Japanese word segmenters

	Analyzer	Speed (sents/s)	Ratio
Rule-based method	JUMAN	8,802	1.00
	MeCab	52,410	0.17
Statistical method	KyTea (Jumandic)	4,892	1.79
	KyTea (Unidic)	1,995	4.41
Juman++: Statistical/Neural hybrid method	V1 noRNN	27	328.82
	V1 RNN	16	535.72
	V2 noRNN	7,422	1.18
	V2 RNN	4,803	1.83

* Results are from [\[Tolmachev+ '20\]](#).

* [Vaporetto](#) and [vibrato](#), implementations of KyTea and MeCab like tokenizers, achieve faster tokenization than them.

Accuracy (F1 scores) of Japanese word segmenters on web text (KWDLC data)



Summary of Part 2

- We have looked at:
 - Progression of lexical analysis methods from statistical to neural models.
 - Examples and characteristics of statistical and neural methods (for Japanese morphological analysis and word segmentation):
 - Optimized statistical models are efficient and fast, whereas neural models are language-independent and have the potential to achieve robust accuracy across domains.
 - Importance of processing speed: practical systems pursue both speed and accuracy.

Part 3: Tokenization in the Neural NLP Era

What are Suitable Token Units for Various NLP Tasks?

● Word Segmentation: *Character*

- Words cannot be used before predicting them.

Example inputs/outputs

Input character tokens:

[奈, 良, に, は, 鹿, が, い, る]



Label sequence to be predicted:

[B, E, B, B, B, B, B, E]

Nara ni-wa shika-ga iru
There are deer in Nara.

B: Beginning of a word
E: End of a word

What are Suitable Token Units for Various NLP Tasks?

- Most other tasks: *Word* (?)

- Using words seems efficient and easy.
- Using characters leads to long token sequences, which increases *computation costs* and raises *modeling difficulty*.

Example inputs/outputs for Sentiment Analysis



When using character tokens:

[I, _, d, o, _, n, o, t, _, l, i, k, e, _, t, h, i, s, _, m, o, v, i, e]

Label to be predicted:
negative



When using word tokens:

[I, do, not, like, this, movie]

A system is required to grasp the information expressed by the combination of these tokens.


* Previous character-level models have shown lower accuracy than (sub)word-level models [\[AI-Rfou+ '19\]](#). However, it has been demonstrated that Transformers with deep layers and a large number of parameters, even at the byte level, can achieve performance competitive with (sub)word-level models [\[Choe+ '19\]](#).

Problems of Using Word Tokens in Downstream Tasks

- A large vocabulary (with size $|V|$) leads to:

- Expensive (infeasible) computation, particularly for language generation models
- A large number of model parameters

Corpus (Set of texts)



Words appeared in the corpus


Model's vocabulary

ID	Token
0	[UNK]
1	the
2	is
...	...
50000	placid
...	...

Neural network model

RNN with attention

(Cited from [Dive into Deep Learning](#))



To keep the vocabulary size practical (typically $< 100K$), a special unknown token is used to represent all words outside the vocabulary, but this is not ideal.

When $|V|=32,000$ and $\text{dim}=512$, word embedding parameters account for 25% of [GRU model](#).

A Solution: Subword

● Subword tokenization

- Breaks down less common words into *subword* tokens based on statistical criteria.

Input text: There have been significant advancements in NLP technologies.

Output tokens:

[“there”, “have”, “been”, “significant”, “advancement”, “##s”, “in”, “nl”, “##p”, “technologies”, “.”]

The symbol “##” represents non-initial token in a word

(tokenizer: google-bert/bert-base-multilingual-uncased)

● Pros:

- Vocabulary size is controllable as a model’s hyper-parameter.
- Unknown words are rare: Most words can be represented by combinations of subwords.

These drawbacks are not often practically critical, and subwords have become the de facto standard.

● Cons:

- Subwords do not align with the boundaries of linguistically meaningful units (e.g., morphemes).
- A tokenizer depends on specific training data, leading to less portability.

* word segmenters also have this limitation.

Subword Tokenization Algorithms

- Three popular algorithms:
 - WordPiece [\[Schuster+ '12\]](#)
 - Byte Pair Encoding (BPE) [\[Sennrich+ '16\]](#)
 - Unigram Language Model (LM) [\[Kudo '18\]](#)
- Two phases for subword tokenization
 - Train a tokenization model from a training corpus.
 - Use the model to tokenize new text.

} Explained in detail in later slides.

cf. https://huggingface.co/docs/transformers/tokenizer_summary.

Byte Pair Encoding (BPE)

● Training algorithm

1. Create an initial small subword vocabulary of all characters in the training corpus.
2. Merge the two consecutive subwords that occur most frequently in the corpus.
 - The tokenizer learns the merge rule.
3. Repeat step 2. until the vocabulary reaches the desired size.

Corpus:

("hug", 10), ("pug", 5), ("pun", 12), ("bun", 4), ("hugs", 5)

Indicates that
"hugs" occurs five times, and so on.

Vocabulary:

["b", "g", "h", "n", "p", "s", "u"]

Merge "u" and "g" (frequency: 20)

Corpus (subword-based):

("h" "u" "g", 10), ("p" "u" "g", 5), ("p" "u" "n", 12), ("b" "u" "n", 4), ("h" "u" "g" "s", 5)

Vocabulary:

["b", "g", "h", "n", "p", "s", "u", "ug"]

New subword

Merge "u" and "n" (frequency: 16)

Corpus (subword-based):

("h" "ug", 10), ("p" "ug", 5), ("p" "u" "n", 12), ("b" "u" "n", 4), ("h" "ug" "s", 5)

...


Unigram Language Model

● Training algorithm

1. Create an initial large subword vocabulary V consisting of a reasonably large number of possible subwords from the training corpus $D = \{x^{(i)}\}$.
2. Calculate the likelihood decrease for each subword $s_m \in V$ when it is removed, and Remove the top k (e.g., 20) percent of subwords that most decrease the likelihood.
 - Single-character subwords are retained regardless of their effect on the loss to avoid the out-of-vocabulary problem.
3. Repeat step 2. until the vocabulary reaches the desired size.

Likelihood decrease or *loss* for s_m


$$\mathcal{L}(D, V) - \mathcal{L}(D, V \setminus \{s_m\})$$



Pre-tokenization-free Tokenizer

● Sentencepiece [\[Kudo+ '18\]](#)

- Subword tokenization tool (\neg algorithm) that implemented BPE and Unigram LM.
- Enabled language-independent tokenization by treating all characters, including whitespace, as usual symbols.



Summary of Part 3

- We have looked at:
 - Suitable token unit: It depends on tasks, but word-level tokenization provides computation efficiency compared to character-level tokenization.
 - A clear drawback of word-level tokenization:
It limits vocabulary size, and thus infrequent words are treated as unknown tokens to avoid impractical computational costs.
 - Subword tokenization as a solution:
This has some (not critical) drawbacks, but has become the de facto standard.
 - Well-known subword tokenization algorithms: BPE and Unigram LM.
 - Sentencepiece: A truly language-independent tokenization tool, particularly useful for unsegmented languages.

Somewhat Focusing on
Word Segmentation

Part 4: The Usage of Foundational Lexical Analysis in Current NLP

How Lexical Analysis is Currently Used: A Case of Japanese NLP

- Direct subword tokenization from raw text performs well in downstream tasks.
- Two-step tokenization is also often adopted.

Named entities: Specific real-world objects/concepts



- [bert-japanese](#)'s tokenizer: Unigram LM
 - _|今日|は|京王|線|で|天氣|の|子|の|聖地|新宿|へ
- [llm-jp-3-1.8b-instruct](#)'s tokenizer:
Unigram LM with MeCab pre-tokenization
 - _|今日は|京王|線|で|天氣|の|子|の|聖地|新宿|へ

Direct subword tokenization

Token boundaries do not align with word boundaries, particularly for named entities.

Two-step tokenization

First perform word segmentation and then build a subword vocabulary based on obtained words.

▶ This reduces cases of boundary conflicts.

* “_” is a symbol representing the beginning of a non-spaced sequence.

* A recommended method of combining Sentencepiece and MeCab is explained [here](#) (in Japanese).

How Downstream Task Accuracy Differ by Tokenizers?

- It depends on tasks, but using words often produces good results (in Japanese NLP).
 - Direct subword tokenization led to a large performance drop for Named Entity Recognition.

Tokenizer		MARC-ja	JSTS	JNLI	JSQuAD	JCQA	NER	UD	Avg.	
	Subword	Morphological	Accuracy	Spearman	Accuracy	F1	Acc	F1	LAS	
Two-step: Use word segmenter as pre-tokenizer	bert-base-japanese	95.5±0.1	85.3±0.3	86.8±0.6	86.4±0.2	76.6±0.8	85.6±0.2	93.3±0.1	87.1	
		Ⓜ MeCab	95.4±0.2	84.2±0.1	88.0±0.4	90.1±0.3	74.1±0.7	83.7±0.8	93.6±0.1	87.0
		Ⓜ Juman++	95.5±0.1	84.6±0.4	87.6±0.4	90.1±0.2	73.8±0.3	85.1±0.6	93.6±0.1	87.2
		Ⓜ Sudachi	95.5±0.1	84.2±0.2	88.2±0.3	90.2±0.2	74.2±0.6	83.5±0.6	93.8±0.1	87.1
		Ⓜ Vaporetto	95.6±0.1	84.8±0.2	87.5±0.3	89.9±0.2	74.2±1.1	84.1±0.9	93.7±0.1	87.1
	BPE (\mathcal{B})	Nothing	95.4±0.2	82.8±0.2	87.2±0.2	88.7±0.3	72.8±0.8	62.9±1.1	93.4±0.1	83.3
		MeCab	95.5±0.1	82.4±0.5	87.5±0.3	89.2±0.3	69.8±0.7	84.0±0.9	93.6±0.1	86.0
		Juman++	95.3±0.3	83.3±0.3	87.7±0.2	89.8±0.3	71.1±0.6	84.7±0.5	93.6±0.1	86.5
		WordPiece	95.3±0.2	83.7±0.3	87.2±0.4	89.6±0.1	70.0±0.9	82.4±0.6	94.0±0.1	86.0
		(\mathcal{W})	Vaporetto	95.3±0.2	83.6±0.1	88.0±0.4	89.7±0.2	71.0±0.4	84.0±0.8	93.8±0.1
	Unigram (\mathcal{U})	Nothing	85.5±0.0	N/A	55.3±0.0	10.1±0.1	20.0±0.8	0.0±0.0	63.8±0.9	33.5
		MeCab	95.4±0.3	84.6±0.4	88.3±0.4	89.5±0.3	74.5±0.8	83.1±1.0	93.4±0.2	87.0
		Juman++	95.4±0.2	84.3±0.3	87.8±0.3	89.9±0.2	74.9±1.2	84.1±0.4	93.4±0.1	87.1
		Sudachi	95.6±0.2	84.8±0.5	88.4±0.3	89.9±0.1	74.5±0.6	83.0±1.3	93.7±0.1	87.1
		Vaporetto	95.5±0.3	84.6±0.2	87.9±0.3	89.9±0.1	74.3±0.8	84.1±0.4	93.7±0.1	87.1
Unigram achieved robust accuracy w/o word segmentation (except for NER)		Nothing	95.4±0.4	83.9±0.3	87.7±0.8	89.3±0.1	74.6±0.4	76.9±1.0	93.2±0.2	85.9

Sentiment analysis Sentence similarity Entailment recognition Question answering Named entity recognition Dependency parsing

* Results are from [Fujii+ '23].

A Case of a Morphologically Rich Language: Kinyarwanda

● KinyaBERT

- Uses two-tier BERT architecture (morpheme and sentence-level encoders) designed specifically for this language.
- Achieved high accuracy across tasks.

Words and morphemes of Kinyarwanda

Word	Morphemes
twagezeyo ‘we arrived there’	tu . a . ger . ye . yo
ndabyizeye ‘I hope so’	n . ra . bi . izer . ye
umwarimu ‘teacher’	u . mu . arimu

Models with BPE or morpheme tokenization yielded lower accuracy than KinyaBERT on most tasks.

Model	MRPC	QNLI	RTE	SST-2	STS-B	WNLI	NER	NEWS
XLM-R	82.6/76.0±0.6/0.6	78.1±0.3	56.4±3.2	76.3±0.4	69.5/68.9±1.0/1.1	63.7±3.9	71.8±1.5	84.0±0.2
BERT _{BPE}	82.8/76.2±0.6/0.8	81.1±0.3	55.6±2.8	79.1±0.4	68.9/67.8±1.8/1.7	63.4±4.1	74.8±0.8	88.3±0.3
BERT _{MORPHO}	82.7/75.4±0.8/1.3	80.8±0.4	56.7±1.0	80.7±0.5	68.9/67.8±1.5/1.3	65.0±0.3	72.8±0.9	86.9±0.3
KinyaBERT _{ADR}	84.4/78.7±0.5/0.6	81.2±0.3	58.1±1.1	80.9±0.5	73.2/72.0±0.4/0.3	65.1±0.0	77.2±1.0	88.0±0.3
KinyaBERT _{ASC}	84.6/78.4±0.2/0.3	82.2±0.6	58.8±0.7	81.4±0.6	74.5/73.5±0.2/0.2	65.0±0.2	76.3±0.5	88.0±0.2

Cited from [Nzeyimana+ '22]

* The results on the test data.

What Tokenization is Used in State-of-the-Art LLMs?


● Byte-level BPE

- The vocabulary is initialized with 256 UTF-8 byte tokens, and new tokens are added by merging existing tokens.
- Completely eliminates the issue of unknown words/characters.

* In UTF-8 encoding, each byte (=8 bits) can represent $2^8=256$ possible values.

Adopted in:

- English-centric LLMs: OpenAI's GPT series, Llama 3, etc.
- Chinese-centric LLMs: Alibaba's Qwen series, etc.



<https://belladoreai.github.io/llama3-tokenizer-js/example-demo/build/>

What Tokenization is Used in State-of-the-Art LLMs?

- Two-step tokenization (word and subword)

- Japanese-centric LLM
 - LLM-jp: MeCab+JumanDIC → Unigram LM

- Japanese-centric LLM extended from English-centric LLM
 - Swallow: MeCab+UniDic → BPE

Efficient tokenization is important to minimize costs.

Input: 日本語の自然言語処理

nihon go-no shizen gengo shori

Tokenizers not optimized for non-Latin characters increase training/inference costs for non-Latin languages.

➤ Before vocabulary expansion (Base model: Llama 2)

→ 日|本|語|の|自|然|言|語|<0xE5>|<0x87>|<0xA6>|理

12 tokens

Byte sequence

➤ After vocabulary expansion (Swallow-7b-hf)

→ 日本|語|の|自然|言語|処理

6 tokens

↳ Japan|word or language|(no-particle)|natural |language|processing

Swallow's vocabulary expansion improved Japanese text generation efficiency up to 78%, while almost maintaining accuracy in downstream tasks.

[Fujii+ '24]

Summary of Part 4

- We have looked at:
 - Variations of tokenization strategies and their performance in downstream tasks:
 - For Japanese, not only direct subword tokenization but also a two-step tokenization (using a word segmenter for pre-tokenization) is common, and the latter provides better accuracy for boundary-sensitive tasks such as NER.
 - For Kinyarwanda, a morphology-aware tokenizer designed for this language achieves high accuracy.
 - Common tokenization approaches in current LLMs:
 - Byte-level BPE eliminates the issue of unknown characters.
 - Efficient tokenization design is critical.

Conclusion: Current and Future of Lexical Analysis Tasks

- Decreased importance in the neural era

- End-to-end neural systems perform well without linguistic information.
- If the goal is to achieve high accuracy in downstream NLP tasks, considering lexical (and higher-level) linguistic analysis is often unnecessary.
 - However, tokenization-related tasks may enhance downstream task accuracy to some extent in unsegmented and morphologically rich languages.

- Unchanging usefulness as fundamental tools

- There is stable demand for constructing new corpora with linguistic information for fields such as (computational) linguistics.
- If the goal is to obtain lexical analysis results, lexical analysis systems are essential.

- Future research for lexical analysis

- Research opportunities remain since the performance of current systems is not perfect.
- Academic research in this area will likely continue, though on a small scale.


* For conducting further research on established tasks, it is important to explain how critical the errors of existing systems are for a target use case.

Assignments

Report Submission

● Specific Instructions for this session

- The deadline is 23:59 on **December 12, 2025**.
- Submit your report in a **plain text file** following the format of **NLP2_report_format.txt**.



```
NLP2_report_format.txt
1,<
2,<
3,<
4,<
5,<
6,<
7,<
8,<
9,<
10,[EOF]
```

● General Instructions

- Submit your reports through <https://edu-portal.naist.jp/>.
- Late submission policy **after the deadline**:
 - We will accept late submission within 3 weeks, but please use the submission portal with Late Submission.
 - Scores will be scaled: 50% when submitted within 1 week, 25% within 2 weeks and 12.5% within 3 weeks.

Assignments

● Instructions

- See "[NLP2_assignments_20251112.ipynb](#)" and provide answers to questions 1-10. } Find the file in the shared folder
I recommend using Google Colab for questions 1-8.
- For questions that require counts, provide only the numerical value.
 - Example: 5
- For questions that choose one options, provide only the option number.
 - Example: a
- Submit your final answers as a single line of text, with answers for Questions 1–10 in order.
Format each answer as "{question number},{answer}". Spaces between "," are optional.
 - Example: 1,5
2,a OR 1, 5
2, a

Make sure to **follow this format** to ensure your answers are processed correctly!

Appendix

Optional Reading Materials

● Tokenization

- [\[Mielke+ 2021\]](#) Between words and characters: A Brief History of Open-Vocabulary Modeling and Tokenization in NLP

● Morphological Analysis

- [\[Liu 2021\]](#) Computational Morphology with Neural Network Approaches
- [\[Baxi+ 2024\]](#) Recent advancements in computational morphology : A comprehensive survey

● Part-of-Speech Tagging

- [\[He+ 2020\]](#) A Survey on Recent Advances in Sequence Labeling from Deep Learning Models
 - Note: This paper includes citations for state-of-the-art POS tagging methods (at the time of publication).
- [\[Chiche+ 2022\]](#) Part of speech tagging: a systematic review of deep learning and machine learning approaches
 - Note: This article mainly reviews journal articles published between 2017-2021 and rarely includes international conference papers, especially those from ACL-related conferences. The coverage of state-of-the-art methods is limited, but it seems useful for understanding a rough trend in POS tagging research.

Optional Reading Materials

● Japanese Morphological Analysis

- [\[Unno 2011\]](#) 形態素解析の過去・現在・未来 (In Japanese)
- [\[Kaji 2013\]](#) 日本語形態素解析とその周辺領域における最近の研究動向 (in Japanese)
- [\[Kudo 2018\]](#) 形態素解析の理論と実装 (Book; In Japanese)
- [\[Higashiyama 2022\] \(slides\)](#) Word Segmentation and Lexical Normalization for Unsegmented Languages
 - Note: Sections 2 and 7 discusses preliminaries and future prospectives for neural word segmentation.

● Chinese Word Segmentation

- [\[Fu+ 2020\]](#) RethinkCWS: Is Chinese Word Segmentation a Solved Task?
- [\[Liu+ 2023\]](#) Survey on Chinese Word Segmentation

● Multilingual Projects for Resource Construction

- UniMorph, <https://unimorph.github.io/>
- Universal Dependencies, <https://universaldependencies.org/>